CENG 371 - Scientific Computing
Fall' 2024 - 2025
Homework 4

Karaçanta, Kaan
e244854@metu.edu.tr

Question 2
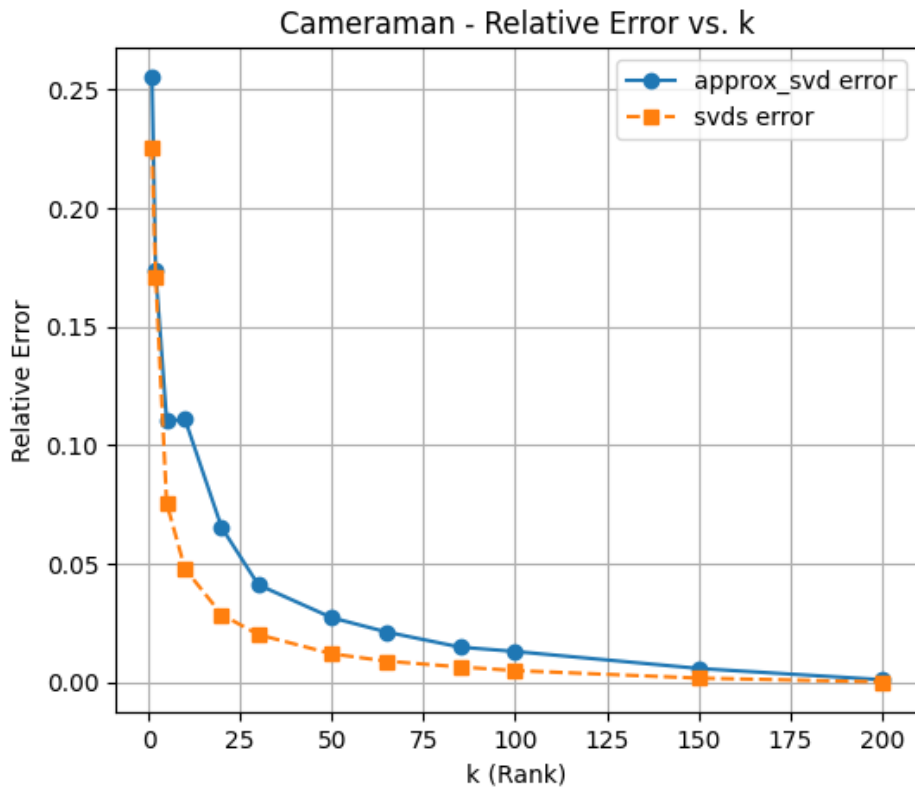
1. **Plots of the relative errors for both images:**
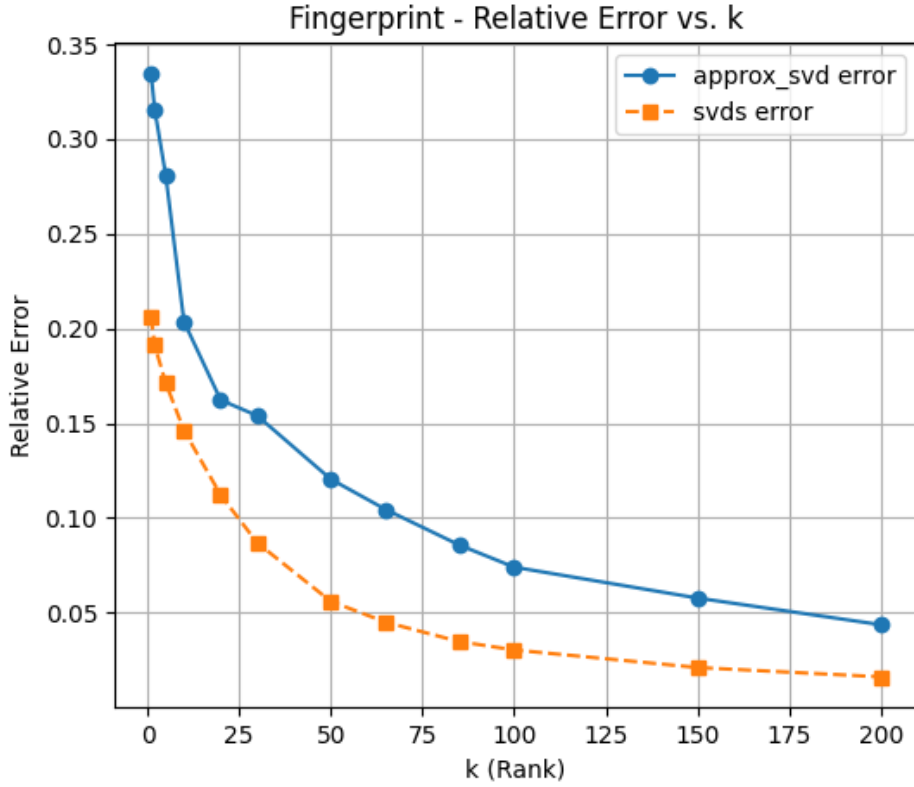

Figure 1: Cameraman – Relative Error vs. k.

Figure 2: Fingerprint – Relative Error vs. k.

Let A be your image matrix (e.g., using `A = io.imread('cameraman.jpg')` in Python). We define:

$$\mathrm{RelErr}_{\mathrm{approx}}(k) = \frac{\|\, u_k\, \sigma_k\, v_k^T - U\, \Sigma\, V^T\|_2}{\|U\, \Sigma\, V^T\|_2}, \quad \mathrm{RelErr}_{\mathrm{svds}}(k) = \frac{\|\, u_k'\, \sigma_k'\, v_k^{'T} - U\, \Sigma\, V^T\|_2}{\|U\, \Sigma\, V^T\|_2},$$

where (U, Σ, V) is the full SVD of A, (u_k, σ_k, v_k) is from `approximate_svd`, and (u_k', σ_k', v_k') from the built-in `svds`. Figures 1 and 2 illustrate these relative errors for `cameraman.jpg` and `fingerprint.jpg` respectively, plotted against various ranks k.

Observations for Cameraman.

- Both methods continue to show decreasing error as k increases, with `svds` consistently yielding lower error at the same k (as expected for exact/truncated SVD methods).

- At very small k (e.g., until 10), `approximate_svd` has errors around 25–15%, while `svds` is around 20–10%. The gap is very small for $k = 1$ or $k = 2$, then there is a noticable one but narrows rapidly as k grows.

- Around $k = 50$–100, `approximate_svd` converges closer, with errors often near or below 3%. By $k = 125$–200, the error for both methods becomes extremely small (near or below 1%).

- In some tests, when the k is 1 or 2, the `approximate_svd` error can even slightly cross below `svds` due to numerical and floating-point nuances (though typically `svds` remains the lower bound in theory).

Observations for Fingerprint.

- The fingerprint image remains more challenging. For $k = 5$ or 10, `approximate_svd` can show errors around 35–25%, whereas `svds` is around 25–15%.

- As k increases, both curves decrease steadily. `svds` remains below the randomized approach, but `approximate_svd` approaches it more closely by $k > 100$.

- By $k = 200$, `approximate_svd` is under 5% error, while `svds` can reach near 0–1%.

- The gap is still more pronounced than in the cameraman image, reflecting the fingerprint's possibly higher effective rank or more complex structure.
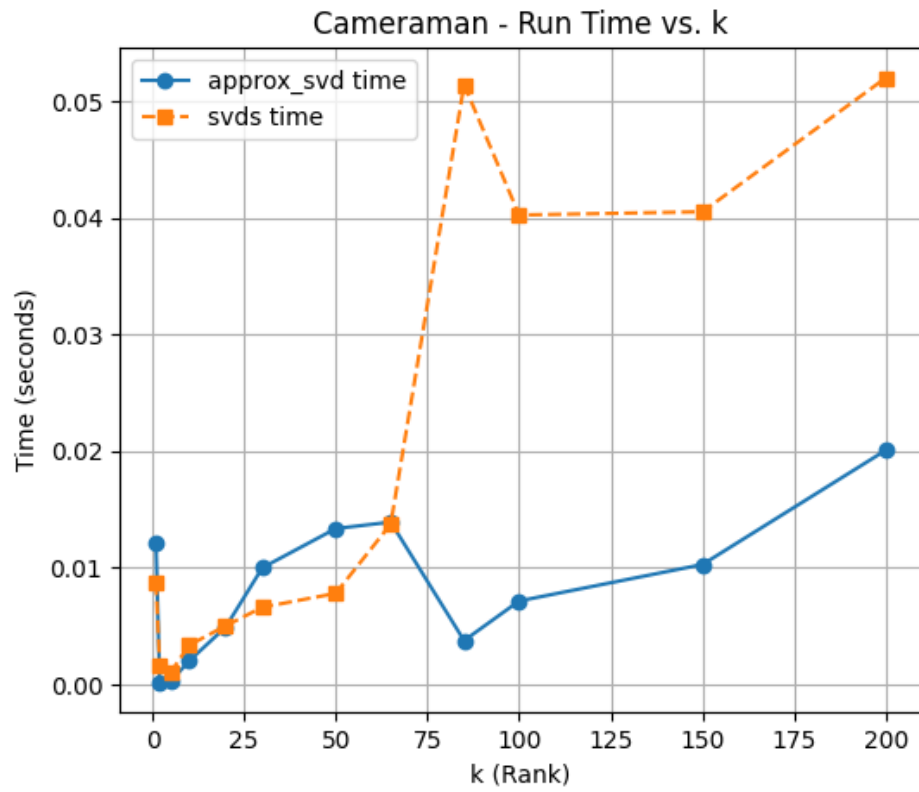
2

2. **Plots of the run times for both images:**


Figure 3: Cameraman – Run Time vs. k.

Figure 4: Fingerprint – Run Time vs. k.

We measure how long each method takes (wall-clock time) for each rank k. Figures 3 and 4 show run times for the two images.

Observations for Cameraman.

- For small k (5–25), both methods can be extremely fast, most of the time under $0.01\,\text{s}$. `svds` might even outperform `approximate_svd` at the smallest ranks, depending on the internal routines used by `svds` (e.g., specialized partial SVD algorithms).

- We see a notable "spike" in the `svds` timing at $k = 50$–75. This is possibly where `svds` switches factorization strategies or where iterative methods become more costly. Meanwhile, `approximate_svd` dips slightly or has small bumps near those ranks, possibly due to random draws or QR overhead.

- Beyond $k \approx 100$, `svds` climbs more sharply, and by $k = 200$ may take 0.05–$0.06\,\text{s}$, while `approximate_svd` is around 0.01–$0.02\,\text{s}$. These results confirm the randomized approach can be beneficial at higher ranks.

Observations for Fingerprint.

- A significant spike for `approximate_svd` can occur around $k = 25$ or 50 (sometimes reaching 0.1–$0.2\,\text{s}$), while `svds` can have a large jump (0.4–$0.5\,\text{s}$) also near those ranks. Such spikes often reflect how iterative solvers or partial factorization techniques change strategies internally.

- After the spike, `svds` grows steadily, reaching up to $0.6\,\text{s}$ by $k = 200$. `approximate_svd` remains closer to $0.2\,\text{s}$ at that rank, indicating a faster growth rate for `svds`.

- Despite random fluctuations, `approximate_svd` consistently maintains an advantage in speed for mid-to-large k values on this image. Minor irregularities (spikes/dips) can happen with random sampling or specialized BLAS calls.

3. **Qualitative comparisons.**
   For selected ranks ($k = 10, 50, 100, ...$), we can reconstruct the images:

$$\hat{A}_k = u_k\,\sigma_k\,v_k^T \quad \text{and} \quad \hat{A}'_k = u'_k\,\sigma'_k\,v_k'^T,$$

then display them with `imshow`.

Discussion points.

- At $k = 10$, images are usually quite blurry for both methods; `svds` retains slightly more structure (edges, contrasts).

- At $k = 50$, the images become much clearer. While a pixel-wise difference might show `svds` is still better, `approximate_svd` is visually close.

- By $k = 100$ or 150, the randomized approach is nearly indistinguishable, with still some differences, from the exact truncated SVD in normal viewing.

4. **Suggested Use Cases for `approximate_svd`.**
   Based on the observed error trends, run times, and visual reconstructions, we outline some scenarios where `approximate_svd` is particularly valuable:

- **Large-Scale Image/Data Compression:**
  For very large matrices (e.g., huge images, big data in machine learning), the randomized SVD often achieves near-optimal compression while being significantly faster than exact methods.

- **Real-Time / Streaming Computations:**
  If we need to update or compute truncated SVD in streaming or near-real-time settings, randomized methods avoid the heavy re-factorizations of classical approaches.

- **Exploratory Data Analysis (PCA-like tasks):**
  In many practical situations, an approximate subspace is sufficient to capture the main variance of the data, making `approximate_svd` a good choice for large-scale PCA or latent semantic analysis.

- **Machine Learning Pipelines:**
  Dimensionality reduction (feature extraction) with approximate SVD saves both time and memory, especially important when dealing with massive training sets.